

CY-4973/7790



Kernel Security: how2rootkit

What is ARM?

- stands for **Advanced RISC Machine(s)**.
- RISC stands for reduced instruction set computer
- ARM is a company, an ambiguous term for Assembly Languages and architectures

ARM Holdings

- A British(?) semiconductor and software company.
- Specializes in designing RISC architectures.
- Does not manufacture its own chips.
- Licenses its designs to other companies (e.g., Qualcomm, Apple, Samsung).
- ARM chips power billions of devices, from smartphones to IoT devices to ssd controllers to enterprise servers.
- Majority owned by Soft Bank (questionable call by the UK)

RISC vs CISC

Feature	RISC	CISC
Instruction Set	Simple, small	Complex, large
Execution	Single cycle per instruction	Multi-cycle per instruction
Power Efficiency	High	Lower
Performance	Optimized for pipelining	Slower due to complexity
Program Size	Larger	Smaller
Examples	ARM, RISC-V	x86, Intel 8086

(Note x86 is kind of like a sugared RISC at this point)

ArmV8 profiles

1. A (Application):

- Supports rich operating systems.
- Focused on performance and complex applications.
- Examples: smartphones, tablets, servers.

2. R (Real-Time):

- Predictable and deterministic performance.
- Common in automotive and industrial systems.

3. M (Microcontroller):

- Low power, low cost.
- Designed for IoT and embedded systems.
- probably in your SSD controller

Microarchitectures

- Specific implementations of the ARM architecture.
- Examples: Cortex-A series, Cortex-R series, Cortex-M series.
- Each microarchitecture optimizes for specific use cases:
 - Performance, power efficiency, or a balance of both.

- This class focuses on ARMv8-a
- RISC arch built by ARM holdings
- aarch64 -> A64: 64 bit execution state
- aarch32 -> A32: 32 bit execution state

A64 (AArch64 mode) Assembly Crash Course

- aarch64: 64 bit execution state for ARMV8-a
- A64 Assembly language/instruction set for aarch64
- 32-bit assembly instructions (4 byte word)
- Uses 31 64 bit general-purpose registers: x0-x30, plus sp (stack pointer) and pc (program counter).
 - w0-w30 are the lower 32 bits b_0, \dots, b_{31}
- Uses 31 128 bit floating point registers q0-q31
 - b,h,w,d,q: ->
 - byte, half-word, word, double-word, quad-word
- Some registers have special usage/convention:
- The goal here is to learn **just enough** to be useful
- technically supports little/big endian but I have only ever seen little endian

Special General Purpose registers (A64)

- Here "special" means there exists a convention
- x0 often used for function return values from a subroutine.
- x0-x7 typically used for arguments in many Linux ABIs.
- x8 holds syscall number
- x29: (sometimes aliased as fp) Frame Pointer (optional)
- x30: Link Register often holds return address on subroutine calls

Comparison to x86_64

Aarch64 Register	Purpose/Usage	x86_64 Equivalent	Explanation
x0	Return values from a subroutine	rax	Both are used to store return values for functions.
x0-x7	Arguments in many Linux ABIs	rdi, rsi, rdx, rcx, r8, r9	x86_64 uses a similar approach with a set of registers for function arguments.
x8	Syscall number	rax	On x86_64, the syscall number is passed through <code>rax</code> before invoking <code>syscall</code> .
x29 (aliased fp)	Frame pointer (optional; helps manage stack frames)	rbp	Both are used as a frame pointer to access local variables and manage the stack.
x30	Link register (holds return address for subroutine calls)	ret stack mechanism	In x86_64, the return address is stored on the stack instead of in a dedicated register.

A64 vs x86_64 notes

- Fixed-length 32-bit instructions simplify decoding and pipelining.
- Implements a **load/store architecture**, meaning it cannot perform operations directly on memory. - Aarch64 cannot directly operate on data unless it is stored in a register.
- All data must first be loaded into registers for processing and then stored back into memory.
 - For example, in `memcpy`, data is copied in chunks by loading blocks into registers and then storing them to the destination address.

Other important registers

- xzr: zero register. ignores writes.
- sp: stack pointer
- pc: instruction pointer/program counter. Can't directly read/modify like x_i
 - bracnh: used to modify pc
- lr: alias for x30.
 - There is no call instruction.
 - There is no ret instruction
- various system registers (exception levels, mmu ...etc)
- this is more trivia than anything but "x31" is either zero register or stack pointer
 - sp and xzr/wzr are **architecturally distinct**.
 - in *some instruction encodings* the same field in the instruction word can represent either sp or xzr depending on context.

Refresher: Common Registers (AArch64)

- **x0-x7:**
- **x8:**
- **x9-x15:** temporary registers.
- **x16-x17:** intra-procedure call scratch regs.
- **x19-x28:** callee-saved registers.
- **x29:**
- **x30:**
- **sp:** stack pointer.
- **pc:** program counter (auto-updated can't be directly used).

Refresher: Common Registers (AArch64)

- **x0-x7:**
 - used to pass function arguments, return values.
- **x8:**
- **x9-x15:** temporary registers.
- **x16-x17:** intra-procedure call scratch regs.
- **x19-x28:** callee-saved registers.
- **x29:**
- **x30:**
- **sp:** stack pointer.
- **pc:** program counter (auto-updated can't be directly used).

Refresher: Common Registers (AArch64)

- **x0-x7:**
 - used to pass function arguments, return values.
- **x8:**
 - indirect syscall param or special usage.
- **x9-x15:** temporary registers.
- **x16-x17:** intra-procedure call scratch regs.
- **x19-x28:** callee-saved registers.
- **x29:**
- **x30:**
- **sp:** stack pointer.
- **pc:** program counter (auto-updated can't be directly used).

Refresher: Common Registers (AArch64)

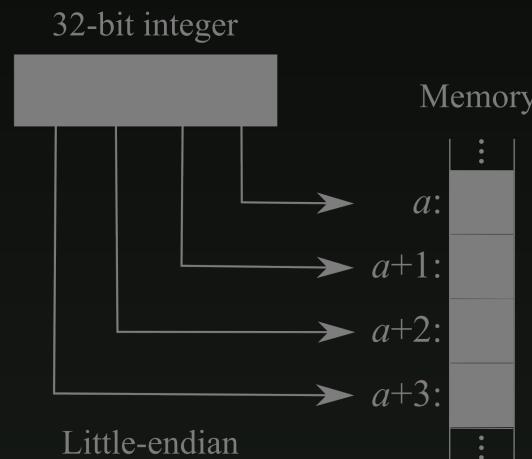
- **x0-x7:**
 - used to pass function arguments, return values.
- **x8:**
 - indirect syscall param or special usage.
- **x9-x15:** temporary registers.
- **x16-x17:** intra-procedure call scratch regs.
- **x19-x28:** callee-saved registers.
- **x29:**
 - optional frame pointer (often used with fp alias)
- **x30:**
- **sp:** stack pointer.
- **pc:** program counter (auto-updated cant be directly used).

Refresher: Common Registers (AArch64)

- **x0-x7:**
 - used to pass function arguments, return values.
- **x8:**
 - indirect syscall param or special usage.
- **x9-x15:** temporary registers.
- **x16-x17:** intra-procedure call scratch regs.
- **x19-x28:** callee-saved registers.
- **x29:**
 - optional frame pointer (often used with fp alias)
- **x30:**
 - link register (return address).
- **sp:** stack pointer.
- **pc:** program counter (auto-updated cant be directly used).

Operand Sizes/ Types

- Byte: 8 bits
- Halfword: 16 bits
- Word: 32 bits
- Doubleword: 64 bits
- Quadword: 128 bits
- ARMv8-a is (almost always) **little-endian** *.
- In assembly, you'll see mnemonics for loading/storing different widths (e.g. `ldrb`, `ldrh`, `ldr`, `ldur`).



Sanity Check

What does this print? #exercise

```
#include <stdio.h>

int main() {
    int x = 0xdeadbeef;
    unsigned char *y = (unsigned char *)&x;
    for (int i = 0; i < 4; i++) {
        unsigned char c = y[i] & 0xff;
        printf("%x ", c);
    }
    printf("\n");
}
```

Sanity Check

What does this print? #exercise

```
#include <stdio.h>

int main() {
    int x = 0xdeadbeef;
    unsigned char *y = (unsigned char *)&x;
    for (int i = 0; i < 4; i++) {
        unsigned char c = y[i] & 0xff;
        printf("%x ", c);
    }
    printf("\n");
}
```

- ef be ad de

Basic Arithmetic (AArch64)

- **ADD / SUB**: Integer addition/subtraction
 - Example: ADD x0, x1, x2 ($x0 \leftarrow x1 + x2$)
 - Immediate form: ADD w3, w4, #10
- **ADDS / SUBS**: Same as above, but **sets condition flags** (N, Z, C, V)
 - Useful for branching on results
- **MUL**: Multiply the lower 64 bits
 - MUL x0, x1, x2
- **SMULH / UMULH**: Signed / unsigned high 64-bit multiply
 - If the product exceeds 64 bits, high part stored in the destination

Load/Store Operations

- **LDR / STR:** Primary load/store instructions
 - LDR $x_0, [x_1] \rightarrow x_0 \leftarrow (\text{uint64_t})x_1$
 - STR $x_2, [x_3, \#16] \rightarrow (\text{uint64_t})(x_3+16) \leftarrow x_2$
- **LDRB / STRB:** For 8-bit (byte) load/store
 - Similarly LDRH, LDRW for halfword/word
- **LDP / STP:** Load/Store **pairs** of registers
 - Often used to save/restore register pairs in function prolog/epilog
- Offsets can be
 - immediate ($[x_1, \#4]$)
 - post-/pre-indexed ($[x_1], \#4, [x_1, \#4]!$)
- READ:
<https://developer.arm.com/documentation/102374/0102/Loads-and-stores---addressing>

Immediate Offset

- The address is computed by adding a constant offset directly to the base register.
- the base register is not modified.

Syntax:

```
ldr w0, [x1, #4] // Load from address (x1 + 4)
str w0, [x1, #-8] // Store to address (x1 - 8)
```

- $\text{address} = \text{x1} + 4$ (or $\text{x1} - 8$).
- Perform the memory access.
- Base register x1 remains unchanged.

Post-indexed Offset

- The address is computed using the base register.
- After the access, the offset is added to the base register.

Syntax:

```
ldr w0, [x1], #4    // Load from x1, then x1 += 4  
str w0, [x1], #-8   // Store to x1, then x1 -= 8
```

- Use the address in the base register for memory access.
- Update the base register ($x1 += 4$ or $x1 -= 8$).

Pre-Index Offset

- The offset is added to the base register before the memory access.
- The base register is updated with the new address.

Syntax:

```
ldr w0, [x1, #4]! // x1 += 4, then load from x1  
str w0, [x1, #-8]! // x1 -= 8, then store to x1
```

- Update the base register ($x1 += 4$ or $x1 -= 8$).
- Use the updated value of the base register for memory access.

Conditional Flags

- Condition Codes: Special flags set by the processor to indicate the result of an operation.
- Enables conditional execution of instructions based on previous operations.
- N (Negative): Set if the result of an operation is negative.
- Z (Zero): Set if the result is zero.
- C (Carry): Set if an operation results in a carry out or borrow.
- V (Overflow): Set if an operation causes a signed overflow.

Instructions-> Conditional Flags

- Comparison Instructions:
 - CMP: Compares two values by subtracting one from the other; updates N, Z, C, V flags.
 - CMN: Compares two values by adding them; updates N, Z, C, V flags.
- Logical Instructions:
 - TST: Performs an AND operation; updates N and Z flags.
 - TEQ: Performs an XOR operation; updates N and Z flags.

Basic Control Flow

- **B** : Unconditional branch
- **BL** : Branch and link (subroutine call)
 - Automatically saves the return address in **x30** (Link Register)
- **RET**: Return from subroutine
 - Not a real instruction :)
 - `blr x30`
 - Jumps to the address in `x30`
- **CBZ / CBNZ**: Compare register to zero and branch
 - `CBZ x0, label` → if `x0 == 0`, branch
- Condition-based branches (after ADDS, SUBS):
 - **B.EQ**, **B.NE**, **B.GT**, **B.LT**, etc.

More Resources

- <https://developer.arm.com/documentation/102374/0102>
- <https://web.archive.org/web/20240829145252/https://modexp.wordpress.com/assembly/>
- <https://mariokartwii.com/armv8/>

How do we learn A64 in this class?

- objdump -d a.out
- Writing little bootstraps

Baby's first A64

```
.section .text
.global _start

_start:
    mov x8, #93          // Syscall number for exit (93)
    mov x0, #0             // Exit code (0)
    svc #0                // Make the syscall
```

Baby's first A64

```
.equ sys_exit, 93

.section .text
.global _start

_start:
    mov x8, sys_exit          // Syscall number for exit (93)
    mov x0, #0                 // Exit code (0)
    svc #0                     // Make the syscall
```

Baby's first A64

```
// Define sys_exit to be 93
.equ sys_exit, 93

// Macro to perform the exit syscall
// note that code is an arg
.macro exit code
    mov x8, sys_exit      // Load syscall number (94 for exit)
    mov x0, \code          // Load the exit code into x0
    svc #0                // Make the syscall
.endm

.section .text
.global _start

_start:
    exit 0                // Call the macro with exit code 0
```

Baby's first A64

```
// common.S
// Define syscall numbers
.equ sys_exit, 93 // Syscall number for exit

// Macro to perform the exit syscall
.macro exit code
    mov x8, sys_exit      // Load syscall number
    mov x0, \code          // Load the exit code into x0
    svc #0                // Make the syscall
.endm
```

```
.include "common.S" // Include the macros file

.section .text
.global _start

_start:
    exit 0
```

hello world

- write: fd=1, addr_str, str_size

```
.section .rodata
hello:
    .asciz "hello world!\n"

.section .text
.global _start
_start:
    // 1 is stdout
    mov x0, #1
    ldr x1, =hello
    mov x2, #14
    mov x8, #64
    svc 0
    mov x0, 0
    mov x8, #93
    svc 0
```

hello world

```
// Macros
.equ STDOUT_FD, 1
.equ SYS_WRITE, 64
.equ SYS_EXIT, 93

.section .rodata
hello:
    .asciz "hello world!\n"

.section .text
.global _start
_start:
    // 1 is stdout
    mov x0, STDOUT_FD
    ldr x1, =hello
    mov x2, #14
    mov x8, SYS_WRITE
    svc 0
    mov x0, 0
    mov x8, SYS_EXIT
    svc 0
```

hello world

```
// Macros
.equ STDOUT_FD, 1
.equ SYS_WRITE, 64
.equ SYS_EXIT, 93

.macro write_stdout message, length
    mov x0, STDOUT_FD
    ldr x1, =\message
    mov x2, \length
    mov x8, SYS_WRITE
    svc 0
.endm

.macro exit_with code
    mov x0, \code
    mov x8, SYS_EXIT
    svc 0
.endm

.section .rodata
hello:
    .asciz "hello world!\n"

.section .text
.global _start
_start:
```

hello world

```
.equ  SYS_EXIT, 93
.section .rodata
fmt: .asciz "Hello %s!\n"
fmt_alt: .asciz "fmt addr: 0x%p\n"
val: .asciz "world"

.section .text
.global _start

_start:
    ldr x0, =fmt
    ldr x1, =val
    bl printf
    ldr x0, =fmt_alt
    ldr x1, =fmt
    bl printf
    mov x8, SYS_EXIT
    svc 0
```

- <https://godbolt.org/z/Gxjor9Kqj>

```
#include <stdio.h>

unsigned long long factorial(unsigned long long n){
    if (n == 0){
        return 1;
    }
    return n * factorial(n-1);
}

int main(int argc, char** argv){
    unsigned long long out = factorial(5);
    printf("Out: %llu\n", out);
return 0;
}
```

```
factorial(unsigned long long):                         // @factorial(unsigned long long)
    cbz    x0, .LBB0_4
    mov    x8, x0
    mov    w0, #1
.LBB0_2:                                         // =>This Inner Loop Header: Depth=1
    subs   x9, x8, #1
    mul    x0, x8, x0
    mov    x8, x9
    b.ne   .LBB0_2
    ret
.LBB0_4:
    mov    w0, #1
    ret
main:                                              // @main
    str   x30, [sp, #-16]!                         // 8-byte Folded Spill
    adrp  x0, .L.str
    add   x0, x0, :lo12:.L.str
    mov   w1, #120
    bl    printf
    mov   w0, wzr
    ldr   x30, [sp], #16                          // 8-byte Folded Reload
    ret
.L.str:
    .asciz  "Out: %llu\n"
```

